

Databricks Certified Data Analyst Associate

Practice Exam — Questions & Answer Key (52 questions)

Correct answers are highlighted in green. Items marked  need a quick check against the screenshot (and two questions weren't captured).

Question 1 verify

A data analyst needs to check if a table called `customer_data` in the `retail.bronze` schema has been tagged with the key `'data_classification'`. Which approach verifies this using Unity Catalog?

- A. `SELECT classification FROM system.information_schema.tables WHERE table_name='customer_data' AND table_catalog='retail' AND table_schema='bronze'`
- B. `SHOW TAGS ON TABLE retail.bronze.customer_data WHERE tag_key='data_classification'`
- C. **`SELECT tag_name, tag_value FROM retail.information_schema.table_tags WHERE table_name='customer_data' AND table_schema='bronze' AND tag_name='data_classification'`**
- D. `SELECT * FROM retail.bronze.customer_data WHERE tag_key='data_classification'`

Answer: C — Query `retail.information_schema.table_tags`.

Why: Tags are exposed via the schema-level `information_schema.table_tags` view, filtered by table/schema/tag name.

Question 2

A data analyst needs to ingest data stored in Amazon S3 into Databricks (batch and streaming). What method should the analyst use?

- A. Download S3 data locally, upload to a UC Volume, then convert
- B. Databricks can read S3 only after files are converted to Parquet
- C. Create a DBFS mount to the S3 bucket via an IAM role, then register the mount as an External Location
- D. **Use Databricks Auto Loader or the `COPY INTO` command to directly ingest files from S3 into Delta Lake tables (supports batch and streaming)**

Answer: D — Use Auto Loader or `COPY INTO` to ingest S3 → Delta directly.

Why: Auto Loader / `COPY INTO` is the supported, governed way to ingest S3 data for both batch and streaming.

Question 2 (alt)

A data analyst cannot read from a table despite the owner running `GRANT SELECT`. Which action should the table owner take to ensure the analyst can access the data?

- A. **Execute `GRANT USAGE (USE CATALOG / USE SCHEMA)` commands at both the catalog and schema levels**
- B. Add the analyst to a group that has `SELECT` privileges on the table
- C. Execute `GRANT USAGE` at the schema level only
- D. Grant `MODIFY` permissions instead of `SELECT`

Answer: A — Grant `USE CATALOG` and `USE SCHEMA` in addition to `SELECT`.

Why: In UC, `SELECT` alone is not enough; the principal also needs `USE` on the parent catalog and schema. (A second 'Question 2' appeared in the set — included as a variant.)

Question 3

A BI analyst applies Data Vault 2.0 to manage historical changes. Which component captures the many-to-many relationships between hubs in a Data Vault v2 model?

- A. Hub Table
- B. Satellite Table
- C. Link Table**
- D. Reference Table

Answer: C — Link Table.

Why: *Links model the many-to-many relationships/associations between hubs.*

Question 4

A data analyst wants to update all records where region is 'East' to set the discount column to 0.1 in a Delta table sales_data. The incorrect attempt used WITH (PARTITIONED BY 'region'). What is the correct query?

- A. MERGE INTO sales_data USING (SELECT 'East' as region) src ON sales_data.region=src.region WHEN MATCHED THEN UPDATE SET discount=0.1;
- B. UPDATE sales_data SET discount = 0.1 WHERE region = 'East';**
- C. UPDATE sales_data SET discount = 0.1 WITH (PARTITIONED BY 'region');
- D. UPDATE sales_data SET discount = 0.1 WHERE region != 'East';

Answer: B — UPDATE sales_data SET discount = 0.1 WHERE region = 'East';

Why: *A simple UPDATE ... WHERE is the correct Delta DML; partitioning clauses don't belong in an UPDATE.*

Question 5

A data analyst wants to compare a Delta table's contents as they existed last week to the current state to track changes. Which Delta Lake feature enables this comparison?

- A. RESTORE command
- B. Time travel querying**
- C. CLONE with VERSION AS OF
- D. VACUUM retention settings

Answer: B — Time travel querying.

Why: *Time travel (VERSION/TIMESTAMP AS OF) lets you read and compare historical table states.*

Question 6

(Screenshot for Question 6 was not captured in extractable text.)

 *Question 6 text wasn't captured from the folder. Share that screenshot and I'll add it.*

Question 7

(Screenshot for Question 7 was not captured in extractable text.)

 *Question 7 text wasn't captured from the folder. Share that screenshot and I'll add it.*

Question 8

A data analyst notices that running the same query multiple times in Databricks SQL returns results much faster after the first execution. Which feature is responsible?

- A. Query Accelerator
- B. Result Cache**
- C. Query Planner
- D. Auto-Optimized Tables

Answer: B — Result Cache.

Why: *Repeated identical queries are served from the result cache, avoiding recomputation.*

Question 9

A data analyst must verify that each nightly upload contains more than a specified number of records and promptly notify IT if not. How can they set up an automated process in Databricks?

- A. Create an alert on the dashboard query that checks the record count, setting a threshold and notifying recipients**
- B. Configure email notifications for all notebook activity regardless of record count
- C. Use dashboard visuals and require IT to manually monitor uploads
- D. Set a static alert in the database table specifying the minimum number of records

Answer: A — Create a Databricks SQL alert on a record-count query with a threshold and notifications.

Why: Databricks SQL alerts evaluate a query against a threshold and notify recipients automatically.

Question 10

How are privileges inherited in Unity Catalog across the hierarchy (metastore, catalog, schema, table/view)?

- A. Privileges granted at a higher level (e.g., a catalog) automatically apply to all child objects unless overridden by more specific permissions at a lower level**
- B. Privileges must be granted individually for each object regardless of hierarchy
- C. Privileges are only inherited from the metastore level; no inheritance between catalogs, schemas, or tables
- D. Privileges granted at the schema level have no impact on tables/views unless directly applied

Answer: A — Higher-level grants flow down to child objects unless overridden.

Why: UC privileges are inherited down the hierarchy; lower-level grants can refine them.

Question 11

How can a data analyst identify poorly performing queries in the Databricks Data Intelligence Platform?

- A. By using Query History and Query Profile tools to examine execution duration, resource usage, and operator-level details**
- B. By using the Table Insights tab to review table usage trends and frequent users
- C. By reviewing raw system logs exported in JSON to manually identify slow queries
- D. By analyzing audit logs and aggregating query execution times

Answer: A — Use Query History + Query Profile.

Why: Query History (with the Query Profile) surfaces durations, resource usage, and operator-level bottlenecks.

Question 12

A BI analyst validating critical_sales_transactions needs to identify upstream tables, contributing notebooks, and any data-quality/sensitivity tags. Which combination of Unity Catalog features should they use?

- A. ANALYZE TABLE ... COMPUTE STATISTICS and query SYSTEM.METASTORE TAGS
- B. Navigate to the table in Catalog Explorer, select the Lineage tab, then review the Overview tab for tags**
- C. Query INFORMATION_SCHEMA directly for dependencies and use DESCRIBE TABLE EXTENDED for tags
- D. Execute SHOW GRANTS to identify data sources and ALTER TABLE ... SET TAGS to view tags

Answer: B — Use Catalog Explorer's Lineage tab and the Overview tab for tags.

Why: Catalog Explorer visualizes lineage (upstream tables/notebooks) and shows tags in the Overview.

Question 13

A data analyst is setting up a Genie space. Which sequence should they use for defining reasonable sample questions to guide users and improve Genie's accuracy?

- A. Create a dedicated notebook per sample question, write SQL in the notebook, and link them via a workspace object reference
- B. Click 'Configure' > 'Context' in the Genie space, use the 'SQL Queries' tab to add sample questions and matching SQL, then review and resolve inconsistencies**
- C. Navigate to 'Configure' > 'Settings' and enter sample questions into the 'Description' field
- D. Upload a spreadsheet of sample questions and manually map each to a SQL query in a separate document

Answer: B — Configure > Context > SQL Queries tab: add sample questions + SQL, then review.

Why: Genie sample questions with example SQL are added in the space's Context/SQL Queries configuration.

Question 14

A data analyst needs to ingest data from an external REST API and store it in a Delta table. Which approach achieves this using API calls?

- A. Use the Databricks SQL REST API to POST raw JSON directly into a Delta table without any processing
- B. Configure Databricks to automatically poll any external API and ingest without code
- C. Use Python's requests library to fetch data from the API, convert the response to a Spark DataFrame, and write it to a Delta table using `.write.format("delta")`**
- D. Upload the API response as a CSV to DBFS and let Databricks auto-convert it to Delta

Answer: C — requests → Spark DataFrame → `write.format("delta")`.

Why: Fetch via requests, build a DataFrame, then write to Delta — the standard API-ingestion pattern.

Question 15

What is the three-level namespace format used to reference objects in Unity Catalog?

- A. catalog-name.table/view/volume-name.schema-name
- B. table/view/volume-name.catalog-name.schema-name
- C. schema-name.catalog-name.table/view/volume-name
- D. catalog-name.schema-name.table/view/volume-name**

Answer: D — catalog-name.schema-name.table/view/volume-name.

Why: UC objects are addressed as catalog.schema.object.

Question 16

A Delta Lake table with billions of records is frequently filtered on `customer_id` and `event_date`. The analyst wants faster filtering without ongoing manual data-layout management. Which action should they take?

- A. Enable liquid clustering using `CLUSTER BY (customer_id, event_date)`**
- B. Use OPTIMIZE with a custom compaction strategy
- C. Add an index on `customer_id` and `event_date`
- D. Configure data skipping statistics for `customer_id` and `event_date`

Answer: A — Enable liquid clustering with `CLUSTER BY (customer_id, event_date)`.

Why: Liquid clustering self-manages layout for the filtered columns without manual partition tuning.

Question 17 verify

A data analyst is designing a Databricks SQL dashboard with interactive filtering via parameters. Which TWO parameter types are supported in Databricks SQL? (Choose 2)

- A. Dropdown List**
- B. Array
- C. String**
- D. Floating Point
- E. Boolean

Answer: A, C — Dropdown List and String.

Why: Databricks SQL parameters include text (String) and dropdown-list types; array/float/boolean aren't dedicated parameter types here.

Question 18

A data analyst wants to use Lakehouse Federation to securely query a SQL Server database. Which SQL statement creates the connection?

- A. CREATE EXTERNAL CONNECTION sqlserver_conn FORMAT SQLSERVER LOCATION '...' CREDENTIALS (...)
- B. CREATE CONNECTION sqlserver_conn TYPE SQLSERVER OPTIONS (host '...', port 1433, database 'mydb', user 'admin', password 'xxxxxx')**
- C. CREATE DATASOURCE sqlserver_conn TYPE SQLSERVER CONFIG (endpoint='...', database='mydb', auth='xxxxxx')
- D. CREATE SERVER sqlserver_conn TYPE SQLSERVER WITH CONNECTION (host='...', database='mydb', authentication='SQL')

Answer: B — CREATE CONNECTION ... TYPE SQLSERVER OPTIONS (host, port, database, user, password).

Why: Lakehouse Federation uses CREATE CONNECTION ... TYPE <db> OPTIONS(...).

Question 19

A data analyst owns prod.sales.orders (used in critical dashboards) and is leaving; the manager must transfer table ownership to the sales_analysts group. Which SQL command should be used?

- A. GRANT OWNERSHIP ON TABLE prod.sales.orders TO sales_analysts;
- B. TRANSFER OWNERSHIP prod.sales.orders FROM current_user TO sales_analysts;
- C. UPDATE TABLE METADATA SET owner='sales_analysts' WHERE table='prod.sales.orders';
- D. ALTER TABLE prod.sales.orders OWNER TO sales_analysts;**

Answer: D — ALTER TABLE prod.sales.orders OWNER TO sales_analysts;

Why: Ownership is transferred with ALTER TABLE ... OWNER TO.

Question 20

A data analyst is setting up a new AI/BI Genie space for natural-language querying with high-quality, governed responses. Which component is essential for an effective Genie space?

- A. Well-annotated datasets registered to Unity Catalog, curated sample queries, and domain-specific instructions to guide natural-language translation**
- B. A dedicated physical server allocated for the Genie space
- C. An external data warehouse configured to store and replicate all data
- D. A standalone chatbot application integrated via custom APIs

Answer: A — Well-annotated UC datasets + curated sample queries + domain instructions.

Why: Genie accuracy comes from good metadata, example queries, and clear domain guidance.

Question 21

An analyst created a single large table containing all product, customer, and sales data instead of a proper star schema, causing slow queries and difficult maintenance. Which action fixes the modeling approach?

- A. Export the data to a spreadsheet for manual filtering, then re-import results
- B. Split the large table into a central fact table for sales transactions and multiple dimension tables (products, customers, time), then connect them using foreign keys**
- C. Merge product, customer, and time info into a single dimension table and join only when necessary
- D. Keep all data in a single table but add more indexes

Answer: B — Split into a central fact table and dimension tables joined by foreign keys.

Why: A proper star schema separates facts from conformed dimensions linked by keys.

Question 22

A data analyst curating a Genie space (curated UC tables, domain instructions, sample questions) wants to maximize trust and reuse of frequently asked queries. Which additional step?

- A. Allow users to edit the underlying datasets
- B. Enable automatic query generation so users get instant responses without review
- C. Disable domain instructions to prevent bias
- D. Mark commonly used queries as Trusted Assets after verifying their correctness**

Answer: D — Mark commonly used queries as Trusted Assets.

Why: Trusted Assets promote verified queries so Genie reuses them reliably.

Question 23

A data analyst needs to sort employee records so that NULL salaries appear last regardless of sort direction. Which SQL statement produces the correct result?

Answer: ORDER BY ... NULLS LAST — Use ORDER BY salary <ASC|DESC> NULLS LAST (or ORDER BY (salary IS NULL), salary ...).

Why: NULLS LAST forces nulls to the end irrespective of ASC/DESC.

⚠ The answer options on this screenshot weren't legible in the captured text. The correct technique is NULLS LAST — share the image and I'll map it to the exact lettered option.

Question 24

A BI analyst needs a chart tracking monthly revenue for three product categories across two years, making it easy to compare each category's trend and spot seasonal patterns. Which visualization type best communicates this?

- A. Multi-series line chart plotting month on the x-axis and revenue on the y-axis for each category**
- B. Single-value counter displaying combined revenue for the two-year period
- C. Scatter chart with revenue on x-axis and profit margin on y-axis
- D. Pie chart showing each category's percentage of total revenue for the final month

Answer: A — Multi-series line chart.

Why: A multi-series line chart shows each category's trend over time for easy comparison.

Question 25

Which SQL clause helps specify the file format for a managed Iceberg table in Databricks SQL?

- A. LOCATION iceberg**

- B. EXTERNAL
- C. USING delta
- D. USING iceberg**

Answer: D — USING iceberg.

Why: The USING clause sets the table format; USING iceberg creates a managed Iceberg table.

Question 26

A data analyst wants to compute the number of orders and the approximate number of distinct customers from a table 'orders' (order_id, customer_id, amount). Which SQL query should be used?

- A. SELECT COUNT(*) AS order_count, customer_id FROM orders;
- B. SELECT COUNT(*) AS order_count, COUNT(customer_id) AS customer_count FROM orders;
- C. SELECT COUNT(*) AS order_count, APPROX_COUNT_DISTINCT(customer_id) AS customer_count FROM orders;**
- D. SELECT COUNT DISTINCT(order_id) order_count, APPROX_COUNT_DISTINCT(customer_id) AS customer_count FROM orders;

Answer: C — COUNT(*) + APPROX_COUNT_DISTINCT(customer_id).

Why: COUNT(*) gives the order count; APPROX_COUNT_DISTINCT gives the approximate distinct customers efficiently.

Question 27 verify

A data analyst shares a Genie space with marketing team members who need to ask questions but should NOT modify its configuration. Which permission levels should be assigned?

- A. CAN MANAGE for the Genie space and CAN USE for the SQL warehouse
- B. CAN VIEW for the Genie space and CAN USE for the SQL warehouse**
- C. CAN EDIT for the Genie space and CAN MODIFY for the SQL warehouse
- D. CAN RUN for the Genie space and CAN VIEW for the SQL warehouse

Answer: B — CAN VIEW (Genie space) + CAN USE (SQL warehouse).

Why: View-level access lets users ask questions without editing; the warehouse needs CAN USE to run queries.

Question 28

A data analyst writing a complex query to find top-performing products by region keeps hitting SQL syntax errors. Which Databricks action helps debug and correct the query?

- A. Enable the 'Auto-complete' feature in the SQL editor
- B. Use Databricks Assistant and click 'Diagnose error' to receive explanations and fix suggestions**
- C. Run EXPLAIN ANALYZE to understand the execution plan
- D. Add a %sql magic command at the beginning of the cell

Answer: B — Use Databricks Assistant → Diagnose error.

Why: The Assistant's Diagnose error explains the failure and suggests corrections.

Question 29

A data analyst needs to reproduce a report as it appeared on July 1, 2025, using a Delta table 'sales_data'. Which SQL statement accesses the data as it existed then?

- A. SELECT * FROM sales_data VERSION AS OF '2025-07-01';
- B. SELECT * FROM sales_data SNAPSHOT AS OF '2025-07-01';
- C. SELECT * FROM sales_data TIMESTAMP AS OF '2025-07-01';**
- D. SELECT * FROM sales_data WHERE date = '2025-07-01';

Answer: C — `SELECT * FROM sales_data TIMESTAMP AS OF '2025-07-01';`

Why: Time travel by date uses `TIMESTAMP AS OF`; `VERSION AS OF` expects a version number.

Question 30

A data analyst joins a local Delta table (`retail.customer_360.customers`) with raw purchase logs in a federated Microsoft SQL Server database (foreign catalog already configured). Which SQL pattern should be used to efficiently join these for exploration?

Answer: Fully-qualified federated join — Join using fully-qualified three-level names, referencing the foreign catalog: `<foreign_catalog>.<schema>.<table>` joined to `retail.customer_360.customers`.

Why: Federated tables are queried by their three-level foreign-catalog name, joined to the local table on the key.

⚠ The answer options on this screenshot were garbled in the captured text. The correct approach is a fully-qualified join across the foreign catalog — share the image and I'll map it to the exact lettered option.

Question 31

A data analyst needs to ingest a new CSV file of sales data and prefers a simple, no-code solution that is secure and integrates with governance tools. Which ingestion method should they use?

- A. Auto Loader
- B. SQL Batch ingestion
- C. UI (Drag-and-drop file ingestion)**
- D. Notebook's scheduled ETL

Answer: C — UI drag-and-drop file ingestion.

Why: For a one-off, no-code CSV load, the add-data UI (drag-and-drop) is the simplest governed option.

Question 32

A data analyst presenting a line chart of total sales over time finds the chart displays incorrectly, shows incomplete data, or the X-axis isn't chronological. Which TWO actions ensure an accurate line chart? (Choose 2)

- A. Check that the X-axis column contains continuous date/time values and is sorted appropriately**
- B. Aggregate the Y-axis (total sales) within the visualization editor if the data is not already aggregated**
- C. Ignore missing data points as Databricks auto-interpolates gaps
- D. Enable 'Auto-format' for the X-axis to auto-interpret date fields
- E. Use a scatter plot instead, as it handles unsorted/missing data

Answer: A, B — Ensure a continuous, sorted date X-axis and aggregate the Y-axis in the viz editor.

Why: A proper sorted date axis plus correct aggregation produces an accurate time-series line chart.

Question 33

A data analyst needs high-quality third-party datasets to enrich customer segmentation, wanting governed, up-to-date, SQL-queryable datasets without manual ingestion. Which Databricks feature is appropriate?

- A. Lakeflow Spark Declarative Pipelines
- B. Databricks Marketplace**
- C. Unity Catalog
- D. Catalog Explorer

Answer: B — Databricks Marketplace.

Why: Marketplace lets you discover and instantly query governed third-party datasets via Delta Sharing.

Question 34

A data analyst reviewing a Genie space notices multiple responses flagged as inaccurate. Which action should they take to improve performance?

- A. Add more tables and columns because increasing data volume improves Genie's answers
- B. Update the instructions and add example SQL queries to guide Genie toward correct, context-aware responses**
- C. Turn off the feedback feature to prevent noise from influencing Genie
- D. Delete the flagged responses to clear biasing prior data

Answer: B — Update instructions and add example SQL queries.

Why: Refining instructions and adding example SQL is the supported way to improve Genie accuracy.

Question 35

A data analyst wants an AI/BI dashboard to automatically refresh every Monday at 6 AM and email results to internal stakeholders. How should they schedule it?

- A. Share the dashboard URL and rely on query result caching to refresh when opened
- B. Attach a SQL warehouse with a scheduled start time and configure an alert on each visualization
- C. Open the Schedule tab in the dashboard, select the refresh time, add subscribers, and publish**
- D. Use the Databricks Jobs UI to create a job that triggers the dashboard refresh and emails

Answer: C — Use the dashboard's Schedule tab: set refresh time, add subscribers, publish.

Why: AI/BI dashboards support scheduled refresh + email subscriptions natively via the Schedule tab.

Question 36

A data analyst works with an IoT dataset where sensor readings are continuously appended to a Delta table, and needs to automatically process new readings in real-time while maintaining a filtered subset of critical sensor data. Which SQL command should be used?

- A. CREATE MATERIALIZED VIEW important_sensors REFRESH ALWAYS AS SELECT ... WHERE device_type='critical';
- B. CREATE STREAMING TABLE important_sensors AS SELECT ... FROM sensor_readings WHERE device_type='critical';**
- C. CREATE TABLE important_sensors USING DELTA TBLPROPERTIES('delta.enableChangeDataFeed'=true) AS SELECT ...;
- D. CREATE VIEW important_sensors AS SELECT ... WHERE device_type='critical';

Answer: B — CREATE STREAMING TABLE important_sensors AS SELECT ...

Why: A streaming table continuously processes newly appended rows in real-time.

Question 37

A data analyst needs to perform complex SQL queries on large datasets quickly and efficiently, without spending time on infrastructure setup or ongoing management. Which Databricks resource should they use?

- A. Jobs Cluster
- B. Scala Notebook
- C. Databricks Delta tables
- D. Databricks SQL Warehouse**

Answer: D — Databricks SQL Warehouse.

Why: SQL warehouses (especially serverless) give managed, high-performance SQL compute with no infra setup.

Question 38

A data analyst sets up a Genie space (sample queries prepared, table/column metadata well-defined). Which action will improve the accuracy and relevance of Genie's responses?

- A. Add only the necessary tables and columns to the Genie space and ensure all metadata in Unity Catalog is well-annotated**
- B. Remove primary/foreign key relationships from Unity Catalog to simplify the model
- C. Include all available tables and columns regardless of relevance
- D. Populate Genie with default sample queries and generic instructions from templates

Answer: A — Add only necessary tables/columns and ensure metadata is well-annotated.

Why: Scoping to relevant, well-annotated objects improves Genie's grounding and accuracy.

Question 39

A data analyst visualizes raw (unaggregated) sales transactions as a bar chart and wants total sales per region directly in the visualization without modifying the SQL. Which approach should they take?

- A. Use the display() function with an aggregation argument in the notebook cell
- B. Add a GROUP BY clause to the SQL query and aggregate sales by region
- C. Export the data to CSV and aggregate in an external tool
- D. Select region as the X-axis and sales as the Y-axis, then set the aggregation type to 'Sum' in the visualization editor**

Answer: D — Set region X-axis, sales Y-axis, aggregation 'Sum' in the viz editor.

Why: The visualization editor can aggregate (Sum) without changing the underlying SQL.

Question 40

What is the scope of the data lineage tracking feature in Databricks Unity Catalog?

- A. Lineage is recorded only for write operations (INSERT/UPDATE/DELETE); reads (SELECT) are not tracked
- B. Lineage is recorded only for UC managed tables, not external tables
- C. Reads and writes to managed tables and registered external locations, across all supported languages, are tracked for lineage**
- D. Lineage is only tracked for queries run in notebooks or the SQL editor, not dashboards or scheduled jobs

Answer: C — Reads + writes to managed tables and registered external locations, all supported languages.

Why: UC lineage covers reads and writes across SQL/Python/Scala/R for managed and external objects.

Question 41

How are analytical models mapped to the Gold layer in the Delta Lake medallion architecture?

- A. The Gold layer stores backup copies of all data for disaster recovery
- B. The Gold layer stores raw, unprocessed data for initial ingestion and schema evolution
- C. The Gold layer is used for data cleaning, deduplication, and joining across domains
- D. The Gold layer contains de-normalized data models optimized for analytics, reporting, and machine learning**

Answer: D — De-normalized models optimized for analytics, reporting, and ML.

Why: Gold holds business-level, de-normalized, consumption-ready models.

Question 42

An analyst's query 'SELECT region, SUM(revenue) ... FROM sales_transactions JOIN customers ...' fails with [TABLE_OR_VIEW_NOT_FOUND] for sales_transactions. The tables are in the analytics schema within the enterprise catalog (UC). What should the analyst do to fix the error?

- A. Prefix both table names with their catalog and schema: enterprise.analytics.sales_transactions and enterprise.analytics.customers**
- B. Escape the table names with backticks
- C. Use the USE SCHEMA command before the query to avoid full qualification
- D. Add only the schema name: analytics.customers and analytics.sales_transactions

Answer: A — Fully qualify both tables as enterprise.analytics.<table>.

Why: With UC, unqualified names fail unless the right catalog/schema context is set; fully qualifying resolves it.

Question 43

A data analyst must share a certified dataset (approved public economic indicators, in UC) with an external research institution that has no workspace access — securely, read-only, only the specified tables, no full-catalog access, with auditability. How should they share it?

- A. Grant the external users access to Unity Catalog by adding their emails to the workspace
- B. Use COPY INTO to export to a managed Delta table, then GRANT SELECT to the institution
- C. Create a Delta Share, add the required table, define a recipient, and have a UC metastore admin/privileged user grant appropriate access**
- D. Use GRANT SELECT to give the external group's principal access to the table

Answer: C — Create a Delta Share with the specific table and a defined recipient.

Why: Delta Sharing securely shares specific tables read-only with external (non-workspace) recipients, governed by UC.

Question 44

An analyst designing a Databricks SQL data model needs fast, flexible analysis of sales by product, store, and time, with easy future extensibility (e.g., promotions, demographics). Which schema design should they recommend?

- A. Star Schema**
- B. Data Vault 2.0 Schema
- C. Normalized (3NF) Schema
- D. Snowflake Schema

Answer: A — Star Schema.

Why: A star schema balances fast query performance, simplicity, and easy extension with new dimensions.

Question 45

A data analyst must share a quarterly dashboard with external stakeholders who are NOT workspace members; they should view it securely without workspace access. Which approach meets the requirements?

- A. Export the dashboard as a PDF and email it
- B. Enable external sharing by publishing the dashboard, turning on embedded credentials, and sharing the link**
- C. Add the stakeholders as workspace admins and grant CAN MANAGE on the dashboard
- D. Assign the stakeholders to the All workspace users group with CAN EDIT permissions

Answer: B — Publish the dashboard with embedded credentials and share the link.

Why: Publishing with embedded credentials lets external viewers see it securely without workspace access.

Question 46

A data analyst cleaning a Unity Catalog table needs to replace all NULL values in the customer_email column with the default 'unknown@company.com'. Which SQL function should they use?

- A. **COALESCE(customer_email, 'unknown@company.com')**
- B. NULLIF(customer_email, 'unknown@company.com')
- C. REPLACE(customer_email, NULL, 'unknown@company.com')
- D. ISNULL(customer_email, 'unknown@company.com')

Answer: A — COALESCE(customer_email, 'unknown@company.com').

Why: COALESCE returns the first non-null value, substituting the default for NULLs.

Question 47

A data analyst team needs production-grade pipelines that define transformations declaratively, automatically enforce data quality expectations, and handle errors without manual intervention. Which Databricks capability addresses these requirements?

- A. Lakeflow Jobs
- B. Lakeflow Connect
- C. Lakeflow Tables
- D. **Lakeflow Spark Declarative Pipelines**

Answer: D — Lakeflow Spark Declarative Pipelines.

Why: Declarative pipelines (LDP/DLT) provide declarative transforms, quality expectations, and managed error handling.

Question 48

Which action should an analyst take to build an AI/BI dashboard and reuse a filter widget across multiple charts on the same dashboard?

- A. Add the widget to a SQL Warehouse configuration so the filter applies globally
- B. **Use the AI/BI Dashboard Canvas to add a filter widget and link it to multiple charts so changing the value dynamically updates all linked queries**
- C. Create a widget in a notebook using dbutils.widgets.dropdown() and reference it in each chart's query
- D. Create a unique filter widget for each chart, since widgets must be scoped individually

Answer: B — Add a filter widget on the AI/BI Dashboard Canvas linked to multiple charts.

Why: Dashboard filter widgets can be linked to many charts so one control updates all of them.

Question 49

What are two specific benefits of implementing Photon as the query execution engine, in terms of performance and cost efficiency? (Choose 2)

- A. Initiates automatic code rewriting for all legacy workloads
- B. **Accelerates Spark SQL and DataFrame query execution**
- C. **Reduces the total cost per workload through faster query performance**
- D. Enables execution of RDD-based transformations at higher speed
- E. Optimizes query performance by automatically tuning cluster configurations

Answer: B, C — Accelerates Spark SQL/DataFrame execution and lowers cost per workload via faster performance.

Why: Photon speeds up SQL/DataFrame workloads, which reduces cost; it doesn't accelerate RDDs or auto-tune clusters.

Question 50

A team built a real-time AI/BI dashboard and wants specific business users/groups to collaborate, with only a few able to make changes while the rest have view-only access — secure, scalable, best-practice. Which approach configures dashboard-level access?

- A. Modify Unity Catalog access on the underlying dataset, granting SELECT to viewers and UPDATE to editors
- B. Request an admin assign workspace permissions, granting Workspace Admin to those who need edit access
- C. Add access-control logic in each widget's SQL query using `current_user()` to restrict by role
- D. Open the dashboard, click 'Share', select individual users or groups, and assign roles such as CAN VIEW, CAN EDIT, or CAN MANAGE**

Answer: D — Use the Share dialog to assign CAN VIEW / CAN EDIT / CAN MANAGE roles.

Why: Dashboard-level sharing with role assignments is the secure, scalable, best-practice approach.

Question 51

What are the core governance capabilities of Unity Catalog?

- A. Automated data pipeline orchestration, ML model deployment, and real-time streaming
- B. Data warehousing optimization, query performance tuning, and cost management across workspaces
- C. Cluster management, notebook versioning, and collaborative development provisioning
- D. Centralized access control with ANSI SQL privileges, automated audit logging, data lineage tracking, and cross-workspace data discovery**

Answer: D — Centralized access control (ANSI SQL privileges), audit logging, lineage, and data discovery.

Why: These are UC's core governance pillars.

Question 52

A data analyst is building a parameterized SQL query to test sales data by region and year with different parameter values before adding it to a dashboard. Which approach allows dynamic testing of the parameters in the Databricks SQL editor?

- A. Hard-code the region and year values and run it for each test
- B. Replace the column names with question marks (?) and expect Databricks to prompt for values
- C. Add the parameters as variables in a Python notebook cell and reference them in the SQL query
- D. Use parameter markers (e.g., :region, :year) in the SQL query and enter their values in the input boxes below the query**

Answer: D — Use parameter markers and enter values in the input boxes below the query.

Why: Named parameter markers create input boxes in the SQL editor for dynamic testing.